Held-out-entity OR discrimination

Model	Scope	R_all	S_all	S_full	Original R_all - S_all	Sensitivity R_all - S_full
Qwen	Pooled	0.983280	0.844643	0.860117	0.138637 [0.118394, 0.155166]	0.123162 [0.101639, 0.141084]
Qwen	Five-topic macro	0.980021	0.828892	0.852042	0.151129 [0.125737, 0.172078]	0.127979 [0.095949, 0.160246]
Llama	Pooled	0.988610	0.811593	0.822829	0.177017 [0.149099, 0.201004]	0.165781 [0.140790, 0.187191]
Llama	Five-topic macro	0.989159	0.831233	0.834045	0.157926 [0.122276, 0.186238]	0.155114 [0.118500, 0.183739]

Bare OR TF/FT versus FF AUROC (TT excluded), averaged at the metric level across samples 11/23/37.
Pooled R_all - S_all is the original primary; macro is a qualification. R_all - S_full is a separate sensitivity.
Brackets are paired pointwise conditional 95% intervals. Rounded values do not determine decisions.