Wording sensitivity of the learned readout

Model	Wording	Sample	Wording - bare [95% interval]
Qwen	Or both	11	-0.079955 [-0.105833, -0.055968]
Qwen	Or both	23	-0.107391 [-0.127501, -0.082662]
Qwen	Or both	37	-0.095878 [-0.119628, -0.071769]
Qwen	At least one	11	-0.083023 [-0.103327, -0.061488]
Qwen	At least one	23	-0.071357 [-0.087614, -0.053585]
Qwen	At least one	37	-0.036443 [-0.047905, -0.024487]
Llama	Or both	11	-0.119558 [-0.139430, -0.097466]
Llama	Or both	23	-0.114614 [-0.133444, -0.093469]
Llama	Or both	37	-0.076054 [-0.093583, -0.057872]
Llama	At least one	11	-0.037791 [-0.053852, -0.022519]
Llama	At least one	23	-0.013147 [-0.022728, -0.004445]
Llama	At least one	37	-0.006782 [-0.015427, 0.002735]

Frozen R_all heads; pooled OR TF/FT versus FF; each sample shown separately. Direction: wording minus bare.
Both OR wording variants were excluded from fitting/selection. No averaging of interval endpoints.
Pointwise conditional 95% intervals. See evidence/extracted-rows.json for exact heads and condition IDs.